

Concept Erasure and Adversarial Recovery in Diffusion Models: A Nudity-Focused Study of ESD and UnlearnDiffAtk

Your name

Course and institution

[Submission date]

Abstract

Text-to-image diffusion models can produce unwanted content, motivating methods that suppress a concept inside the model rather than only filtering its final output. This report studies nudity erasure by Erased Stable Diffusion (ESD) and adversarial recovery by UnlearnDiffAtk, also called a diffusion-model unlearning attack. We explain the defense, the white-box prompt attack, and a paired evaluation design. In a matched 20-prompt run, broad NudeNet detections fell from 20/20 for original Stable Diffusion to 7/20 for ESD, then rose to 17/20 after attack. Under a stricter exposed-region rule, detections changed from 11/20 to 1/20 and then 5/20. These results show both substantial erasure and partial adversarial recovery under the stated detector and sample; they do not prove that every detection is explicit nudity or that unlearning fails universally. Every example image in this report is a separate copy censored using padded NudeNet boxes.

1 Introduction

Text-to-image diffusion models are useful because a short prompt can produce varied, realistic images. The same ability creates a safety problem when a user requests nudity or finds a prompt that elicits it indirectly. A post-generation filter can hide an output, but it does not change the image-generating model. Concept erasure instead modifies the model so that the unwanted visual concept is less likely to be generated in the first place [1].

An erasure method should be judged in two ways. First, does ordinary prompting produce fewer detections after editing? Second, does the edited model remain robust when an attacker actively searches for a prompt that restores the erased behavior? The second question is important in information security: a model that appears safe under straightforward tests may still expose the target concept under adaptive inputs [2].

This report restricts the target concept to nudity, the defense to ESD, and the attack to UnlearnDiffAtk. The central question is whether the ESD-edited model suppresses nudity under the original prompts and whether optimized prompts can recover it. Here “erasure” means an observed change in generated behavior, not a guarantee that a concept or training record has been removed from the weights. All generated examples included below are separate report copies masked using NudeNet bounding boxes; quantitative evaluation uses the raw images.

2 Background and threat model

2.1 Text-conditioned latent diffusion

Stable Diffusion encodes images into a latent space and generates by iteratively denoising a random latent. At time step t , a U-Net predicts noise $\epsilon_\theta(z_t, c, t)$ from noisy latent z_t and text condition c . Classifier-free guidance combines conditional and unconditional predictions




Figure 1: Evaluation flow. The original model and the ESD checkpoint must be sampled with matched prompts and generation settings. The attack changes only the prompt supplied to the ESD model. Score raw outputs before preparing masked report copies.

to strengthen the effect of a prompt. Both methods in this report work with this denoising mechanism: ESD changes its weights, while the attack changes the text condition sent to the edited model.

2.2 Attacker and evaluation scope

The attacker has white-box access to the ESD checkpoint and may prepend five optimized tokens to a fixed original prompt. The attacker cannot change the victim weights or the evaluation rule. A target reference image is available for the attack objective. This is a stronger setting than merely trying a few manually written prompts, but it does not represent every possible adversary. We count an output as detected when NudeNet reports a selected class above threshold 0.45. Detection is an operational proxy, not a human judgment of severity.

3 Erasure method: Erased Stable Diffusion

ESD fine-tunes a pre-trained diffusion model to imitate a target noise prediction that points away from an unwanted concept [1]. Let c_e denote the concept to erase, here nudity. With a frozen original model θ^* , a simplified form of the target is

$$\epsilon_{\text{target}}(z_t, c_e, t) = \epsilon_{\theta^*}(z_t, \emptyset, t) - \eta[\epsilon_{\theta^*}(z_t, c_e, t) - \epsilon_{\theta^*}(z_t, \emptyset, t)], \quad (1)$$

where η controls erasure strength and $\emptyset$ is the empty condition. The edited U-Net is trained to match this target. Intuitively, the edit subtracts the model’s learned direction associated with the concept. Unlike an inference-time negative prompt, the change is stored in model weights.

The ESD paper distinguishes ESD-x, which updates cross-attention weights, from ESD-u, which updates non-cross-attention weights. The authors favor ESD-u for broad nudity suppression because the target concept may appear even when a prompt never names it [1]. Our repository loads a checkpoint named `Nudity-ESDx1-UNET-SD.pt`. Its name alone does not establish the exact training variant, so the experiments below refer to it as the supplied ESD nudity checkpoint until provenance is verified. A stronger edit may also alter benign images; utility therefore belongs in a full defense assessment, although it is not measured in the present local run.

4 Attack method: UnlearnDiffAtk

UnlearnDiffAtk searches for a text perturbation that makes an unlearned diffusion model generate a target concept again [2]. For a target image latent z_0^* , the method adds noise at sampled times and searches for an adversarial prompt c_{adv} that lowers the edited model’s denoising error:

$$\min_{c_{\text{adv}}} \mathbb{E}_{t, \epsilon} \left[\|\epsilon - \epsilon_{\theta_{\text{ESD}}}(z_t^*, c_{\text{adv}}, t)\|_2^2 \right]. \quad (2)$$

Table 1: Local setup and items to verify before final submission.

Component	Setting
Base model	CompVis/stable-diffusion-v1-4
Erased model	Supplied Nudity-ESDx1-UNET-SD.pt U-Net; training variant and hash to verify
Prompt sample	Seeded 20-prompt subset; sample seed 2024, run seed 0
Generation defaults	512 $\times$ 512 pixels, 50 sampling steps, guidance 7.5; verify saved per-case overrides
Attack	Five prefix tokens; 50 sampled times; 40 optimization steps/time; Adam, learning rate 0.01
Detector	NudeNet, score > 0.45 ; broad and stricter class rules defined below
Report artifacts	20 matched images per condition; padded NudeNet-masked copies for publication

This objective uses the victim model itself rather than a NudeNet gradient. A prompt whose denoising trajectory better fits the target latent is then tested by actual image generation. NudeNet judges the generated output afterward.

In the local configuration, the attacker optimizes five prefix-token positions. The implementation uses a projected token distribution with Adam, an L2 denoising loss, 50 sampled diffusion times, up to 40 optimization steps per sampled time, learning rate 0.01, and seed 0. The attack iterates until its success check passes or its configured search is exhausted. These are attack-budget details, not properties of ESD itself. Figure 1 separates the weight-edit stage from the prompt-optimization stage.

5 Experimental design

5.1 Paired comparisons

The intended experiment has three conditions for each prompt p_i and generation seed s_i :

1. Original Stable Diffusion v1.4 with p_i (unedited baseline).
2. The supplied ESD checkpoint with the same p_i (erasure baseline).
3. The same ESD checkpoint with an optimized prefix plus p_i (attack output).

Comparing conditions 1 and 2 measures erasure; comparing 2 and 3 measures attack recovery. Prompt, seed, image size, sampler, sampling steps, guidance, and detector must be aligned across conditions. The reference images used by the attacker must *not* be substituted for condition 1, because they are attack targets, not necessarily outputs from the matched original model.

The existing random subset contains 20 prompts, sampled with seed 2024 and evaluated with run seed 0. A separate 20-prompt stress set is useful as a secondary check but should not be pooled with this random set. The model sampler’s defaults are 512-pixel square images, 50 inference steps, and guidance 7.5; the final manifest should confirm any per-prompt seed or guidance overrides. Table 1 records the reproducibility details currently known.

5.2 Metric and image safety

For n prompts, define the detection rate as $n^{-1} \sum_{i=1}^n \mathbb{1}[\text{a selected NudeNet class for image } i \text{ scores } > 0.45]$. The repository’s *broad* rule includes exposed feet, belly, armpits, and male chest in addition to exposed breasts, genitals, buttocks, and anus. This rule should be reported exactly




Figure 2: Matched erasure examples: original SD v1.4 versus ESD for identical prompts, seeds, and generation settings. Case 0 changes from a strict detection to no detection; case 1 retains a broad, non-strict detection. Images are padded NudeNet-masked report copies; scores use the unmasked outputs.

because it determined attack success. We also show a *stricter* descriptive rule limited to exposed female breast, female or male genitalia, buttocks, and anus. Neither metric alone proves what a human viewer would consider explicit nudity.

Detector scoring must use the unmasked generated images. For publication, create separate copies with opaque, padded masks over every NudeNet exposed-region box, then manually inspect for missed sensitive regions. If coverage is uncertain, censor the entire person or omit the example. Apply this rule to *every* figure column, including any original-model baseline. The masks are a presentation safeguard and must not change the quantitative scores.

6 Results

6.1 Erasure under ordinary prompts

With matched prompts, seeds, and generation settings, broad detections fell from 20/20 for original Stable Diffusion v1.4 to 7/20 for ESD, a reduction of 13 images or 65 percentage points. Under the stricter rule, detections fell from 11/20 to 1/20, a reduction of 10 images or 50 percentage points. Figure 2 shows one successful erasure (case 0) and one residual broad-rule detection after ESD (case 1). These examples are selected to display both outcomes rather than only the strongest visual result.

6.2 Attack recovery on the ESD checkpoint

In the available seeded random 20-prompt run, the repository’s broad NudeNet rule detected 7/20 ESD images before attack and 17/20 after attack. Ten of the 13 baseline failures became detected, or 10/13 (about 77%). These paired counts are stronger evidence of recovery than the difference between the two percentages alone. Under the stricter exposed-region rule, final-image detections rose from 1/20 to 5/20, with four new detections. This stricter analysis is post hoc: the attack optimized and stopped under the broad rule. The separate stress set recorded 2/20 before and 16/20 after attack under the broad rule, with 14 new detections among 18 earlier failures. Because its prompts were selected differently, it is a secondary observation, not an additional 20 random trials.

Table 2: NudeNet outcomes. “New” counts previously undetected ESD prompts that became detected after attack.

Set and detection rule	Original SD	ESD before	ESD after	New / prior failures
Random 20, broad	20/20	7/20	17/20	10/13
Random 20, stricter (post hoc)	11/20	1/20	5/20	4/19
Stress 20, broad	—	2/20	16/20	14/18




Figure 3: Qualitative examples from the random 20-prompt run. Columns show original SD, ESD with the original prompt, and the same ESD checkpoint after adversarial-prefix optimization. Opaque masks cover exposed-region boxes returned by NudeNet on separate report copies. Detector decisions and scores were computed on the unmasked images.

Figure 3 compares all three conditions. Case 0 is a broad-rule recovery (**FEET_EXPOSED**, 0.659), case 2 is a stricter-rule recovery (maximum score 0.803 for **FEMALE_BREAST_EXPOSED**), and case 7 is an attack failure with no broad-rule detection. These rows were selected by outcome category rather than visual appeal. Figure 4 summarizes the broad and strict rates for the matched random set.

7 Discussion and limitations

The matched experiment shows that ESD substantially reduced NudeNet detections relative to the original model, while optimized prompts recovered much of the broad-rule detection rate. It does not show that all 17 detected attacked outputs contain explicit nudity: the broad rule counts several exposed but often non-explicit body regions. The stricter 5/20 result is a useful sensitivity check, but it was calculated after attack search under a different success condition.

Several factors limit generalization. Twenty prompts give coarse five-percentage-point increments; the stress set is selected rather than random; stochastic sampling can change individual images; and NudeNet may produce false positives or negatives. A stronger evaluation would

Original SD (broad)		20/20 (100%)
Original SD (strict)		11/20 (55%)
ESD before attack (broad)		7/20 (35%)
ESD before attack (strict)		1/20 (5%)
ESD after attack (broad)		17/20 (85%)
ESD after attack (strict)		5/20 (25%)

Figure 4: Detection counts for the matched random 20-prompt set under broad and stricter NudeNet rules. ESD sharply reduces both rates, while the adversarial prefix recovers a substantial part of the broad-rule rate.

repeat seeds, audit detector labels by human review under a safe procedure, and measure image quality or prompt adherence. The result is best interpreted as a security warning about one checkpoint and one attack budget, not a universal failure rate for all unlearning methods.

8 Conclusion

ESD attempts to suppress a concept by changing diffusion-model weights. UnlearnDiffAtk probes whether an optimized prompt can still recover the target behavior. In the matched local run, broad detections changed from 20/20 on original Stable Diffusion to 7/20 on ESD and 17/20 after attack; the stricter counts were 11/20, 1/20, and 5/20. Thus ESD was effective under ordinary prompting, but the white-box attack partially restored the erased behavior. All illustrative outputs were scored before censoring and are shown only as masked copies.

Remaining work before submission

Verify the checkpoint source and record the NudeNet package version. Manually inspect the masked report copies for missed sensitive regions, replace the author and course placeholders, and adapt the file to the required course template. The expensive generation and attack runs need not be repeated unless the protocol changes.

References

- [1] R. Gandikota, J. Materzynska, J. Fiotto-Kaufman, and D. Bau, “Erasing Concepts from Diffusion Models,” *ICCV*, 2023. <https://arxiv.org/abs/2303.07345>.
- [2] Y. Zhang et al., “To Generate or Not? Safety-Driven Unlearned Diffusion Models Are Still Easy To Generate Unsafe Images ... For Now,” *ECCV*, 2024. <https://arxiv.org/abs/2310.11868>.
- [3] NudeNet, project repository. <https://github.com/notAI-tech/NudeNet>.